

COM SCI 118 Winter 2018

Computer Network Fundamentals

Project 1: Web Server Implementation using BSD Sockets
Due date: Feb. 2nd, 11:59 p.m. PST

1 Goal

In this project, we are going to develop a Web server in C/C++. Also, we will take the chance to learn a little bit more about how the Web browser and server work behind the scene.

2 Lab Working Environment

You need a plain-text editor to create the source code. You can use vim/emacs/nano in Unix/Linux systems. Since C/C++ is a cross-platform language, you should be able to compile and run your code on any machine with C/C++ compiler and BSD socket library installed.

You should test and develop your code in Ubuntu 16.04.3 LTS (Xenial Xerus), as your submission will be tested in this environment. You can obtain a copy of Ubuntu 16.04.3 from its official website (<http://releases.ubuntu.com/16.04/>, 64-bit PC (AMD64) desktop image) and install in VirtualBox (<https://www.virtualbox.org>) or other similar VM platform. We also provide pre-installed virtual machine images on our server (username/password: cs118/cs118):

- For VirtualBox: http://metro.cs.ucla.edu/cs118/CS118_OVF.rar
- For VMware: <http://metro.cs.ucla.edu/cs118/CS118.rar>

It is recommended to work your project in Linux/Unix/macOS systems. It is also highly recommended that testing your code in the specified environment. Windows workstations are NOT recommended, because they have their own idiosyncracies which may be too much for this simple project. Also, TAs will not be able to help you on any of these platforms.

3 Instructions

1. Read Chapter 2 of the textbook carefully. The textbook will help you understand how HTTP works. You can also use the socket programming materials covered by your TA to learn more on the C/C++ socket programming.

You must program in C/C++ (not Java). To help you get familiar with socket programming, we provide a demo client-server code at <https://wing1.cs.ucla.edu/gitlab/zyuan/cs118-project1-demo> (public repo, no need to register for an account). This code is also included in the VM we provide above.

2. Part A: Implement a “Web Server” that dumps HTTP request messages to the console. This is a good chance to observe how HTTP works. You should first start your web server, and

then initiate a web client. For example, you may open Mozilla Firefox and connect to your web server. Your web server should print out the HTTP request message it received. You should be able to explain what do the fields in the message mean by looking up in the textbook or RFC 1945.

3. Part B: Complete your “Web Server” by responding to client’s HTTP request. The “Web Server” should parse the HTTP request from the client, creates an HTTP response message consisting of the requested file preceded by header lines, then sends the response directly to the client.

If the requested file does not exist, your server should be able to display an HTML page showing 404 not found error. Your server does not need to handle requests for files in any subdirectories. All files are supposed to be accessed in the same level as the actual server program. However, your server does need to handle the case where file name contains space.

Your server should at least support HTML files (i.e. *.html and *.htm). It should also support for JPEG and GIF images (i.e. *.jpg, *.jpeg, and *.gif). All the file extensions are case insensitive.

4. Pay attention to the following issues when you are implementing and testing the project.

If you are running the browser and server on the same machine, you may use localhost or 127.0.0.1 as the name of the machine. Instead of using port 80 or 8080 or 6789 for the listening socket, you should pick your own to avoid conflicts. It is suggested not to use port numbers 0–1024 (these are reserved ports).

After you are done with both parts, you need to test your server. You can first put an HTML file in the directory of your server program. Then, you should connect to the server from a browser using the URL `http://<machinename>:<port>/<filename>` and see if it works. For your server in Part A, you should be able to see HTTP requests in the console of your server machine. For your server in Part B, your browser should be able to show the content of the requested file (or displaying image).

4 Grading Criteria

Your code will be first checked by a software plagiarism detecting tool. If we find any plagiarism, you will not get any credit and we are obliged to report to the Dean of Student. Your code will be graded based on several testing rubrics. We list the main grading criteria below; more details will be announced via CCLE.

- The server program compiles and runs normally.
- The server can print out the correct HTTP request messages to the console.
- The server program transmits a small binary file (up to 512 bytes) correctly.
- The server program transmits a large binary file (up to 100 MiB) correctly.
- The server program serves an HTML file correctly and it can show in the browser.
- The server program serves an image file correctly and it can show in the browser.
- The server program serves a file which file name contains space correctly in the browser.

5 Project Submission

Project due date is 11:59 p.m. on Friday, Feb. 2nd on CCLE. Late submission is allowed by Sunday, Feb 4th (10% deduction on Saturday, 20% deduction on Sunday). Put all your files into a directory, compress the directory and generate a “UID.tar.gz” where UID is your UCLA ID. Only one member of your team is required to submit your project to CCLE. Your submission should include the following files:

- Your source codes (e.g. webserver.c). The code of the server can be more than one file.
- A Makefile. The TAs will only type “make” to compile your code.
- Your report in PDF format only (report.pdf).
- A README file which will contain student names and student IDs.

Your report should not exceed 6 pages. In your report:

1. Give a high-level description of your servers design.
2. What difficulties did you face and how did you solve them?
3. Include a brief manual in the report to explain how to compile and run your source code (if TAs cant compile and run your source code by reading your manual, the project is considered not to be finished).
4. Include and briefly explain some sample outputs of your client-server (e.g. in Part A you should be able to see an HTTP request). You do not have to write a lot of details about the code, but just adequately comment your source code. The cover page of the report should include the name of the course and project, your partners name, and student id.